

SECTION A

INSTALL HUGGING FACE & GENERATE TEXT

In [1]: `!pip install transformers`

Requirement already satisfied: transformers in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (5.2.0)

Requirement already satisfied: huggingface-hub<2.0,>=1.3.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (1.4.1)

Requirement already satisfied: numpy>=1.17 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (2.4.2)

Requirement already satisfied: packaging>=20.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (25.0)

Requirement already satisfied: pyyaml>=5.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (6.0.3)

Requirement already satisfied: regex!=2019.12.17 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (2024.9.11)

Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (0.22.2)

Requirement already satisfied: typer-slim in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (0.24.0)

Requirement already satisfied: safetensors>=0.4.3 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (0.7.0)

Requirement already satisfied: tqdm>=4.27 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from transformers) (4.67.2)

Requirement already satisfied: filelock in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from huggingface-hub<2.0,>=1.3.0->transformers) (3.20.3)

Requirement already satisfied: fsspec>=2023.5.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from huggingface-hub<2.0,>=1.3.0->transformers) (2025.10.0)

Requirement already satisfied: hf-xet<2.0.0,>=1.2.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from huggingface-hub<2.0,>=1.3.0->transformers) (1.2.0)

Requirement already satisfied: httpx<1,>=0.23.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from huggingface-hub<2.0,>=1.3.0->transformers) (0.28.1)

Requirement already satisfied: shellingham in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from huggingface-hub<2.0,>=1.3.0->transformers) (1.5.4)

Requirement already satisfied: typing-extensions>=4.1.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from huggingface-hub<2.0,>=1.3.0->transformers) (4.15.0)

Requirement already satisfied: anyio in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from httpx<1,>=0.23.0->huggingface-hub<2.0,>=1.3.0->transformers) (4.12.1)

Requirement already satisfied: certifi in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from httpx<1,>=0.23.0->huggingface-hub<2.0,>=1.3.0->transformers) (2026.1.4)

Requirement already satisfied: httpcore==1.* in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from httpx<1,>=0.23.0->huggingface-hub<2.0,>=1.3.0->transformers) (1.0.9)

Requirement already satisfied: idna in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from httpx<1,>=0.23.0->huggingface-hub<2.0,>=1.3.0->transformers) (3.11)

Requirement already satisfied: h11>=0.16 in /Library/Frameworks/Python.fra

mework/Versions/3.12/lib/python3.12/site-packages (from httpcore==1.*->httpx<1,>=0.23.0->huggingface-hub<2.0,>=1.3.0->transformers) (0.16.0)
 Requirement already satisfied: typer>=0.24.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from typer-slim->transformers) (0.24.0)
 Requirement already satisfied: click>=8.2.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from typer>=0.24.0->typer-slim->transformers) (8.3.1)
 Requirement already satisfied: rich>=12.3.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from typer>=0.24.0->typer-slim->transformers) (14.3.2)
 Requirement already satisfied: annotated-doc>=0.0.2 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from typer>=0.24.0->typer-slim->transformers) (0.0.4)
 Requirement already satisfied: markdown-it-py>=2.2.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from rich>=12.3.0->typer>=0.24.0->typer-slim->transformers) (4.0.0)
 Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from rich>=12.3.0->typer>=0.24.0->typer-slim->transformers) (2.19.2)
 Requirement already satisfied: mdurl~=0.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from markdown-it-py>=2.2.0->rich>=12.3.0->typer>=0.24.0->typer-slim->transformers) (0.1.2)

[notice] A new release of pip is available: 26.0 -> 26.0.1

[notice] To update, run: `pip install --upgrade pip`

```
In [2]: from transformers import (pipeline # Load small model for text generation
      , AutoModelForSeq2SeqLM # Load models for summarization
      , AutoTokenizer # Load models and tokenizers for generation
      )

from IPython.display import display, Markdown
from warnings import filterwarnings
filterwarnings('ignore')
import gensim.downloader as api
import pandas as pd
import urllib.request, zipfile, os, ssl
from gensim.models import KeyedVectors

import numpy as np
import matplotlib.pyplot as plt
import matplotlib.colors as mcolors
import seaborn as sns
from sklearn.metrics.pairwise import cosine_similarity
```

```
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/tqdm/auto.py:21: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
```

```
In [3]: generator = pipeline('text-generation', model='distilgpt2')
```

```
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.
Loading weights: 100%|██████████| 76/76 [00:00<00:00, 1670.21it/s, Materializing param=transformer.wte.weight]
```

```
In [4]: def generate_text(prompt, max_length=50):
      """
```

```

Generate text using the pre-trained model
- can be modified to use a different model or parameters as needed
- generator currently uses 'distilgpt2' which is a smaller, faster ve
- generator can be passed as a parameter to allow flexibility in which
"""
generated = generator(prompt, max_length=max_length, num_return_sequences=1)
return generated[0]['generated_text']

```

```

In [5]: # Generate text based on a prompt
prompts = [
    "Once upon a time... ",
    "AI is transforming industries by... ",
    "OpenClaw is what to AI? "
]

## loop through prompts, generate text, and store results in a DataFrame
rows = []
for prompt in prompts:
    response = generate_text(prompt, max_length=30)
    rows.append({
        "prompt": prompt,
        "model": "distilgpt2",
        "generated_text": response
    })
df = pd.DataFrame(rows)
display(df)

```

Passing `generation\_config` together with generation-related arguments= ({'num_return_sequences', 'max_length'}) is deprecated and will be removed in future versions. Please pass either a `generation\_config` object OR all generation parameters explicitly, but not both.

Setting `pad\_token\_id` to `eos\_token\_id`:50256 for open-end generation.

Both `max\_new\_tokens` (=256) and `max\_length` (=30) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Setting `pad\_token\_id` to `eos\_token\_id`:50256 for open-end generation.

Both `max\_new\_tokens` (=256) and `max\_length` (=30) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Setting `pad\_token\_id` to `eos\_token\_id`:50256 for open-end generation.

Both `max\_new\_tokens` (=256) and `max\_length` (=30) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

	prompt	model	generated_text
0	Once upon a time...	distilgpt2	Once upon a time... _____ \nPosted...
1	AI is transforming industries by...	distilgpt2	AI is transforming industries by... A very p...
2	OpenClaw is what to AI?	distilgpt2	OpenClaw is what to AI? \n\nThe AI is an inte...

TOKENISATION USING HUGGING FACE

```
In [6]: tokenizer1 = AutoTokenizer.from_pretrained('distilgpt2')
tokenizer2 = AutoTokenizer.from_pretrained('bert-base-uncased')
sentence = "LLMs are powerful tools for natural language understanding."
```

```
In [7]: def tokenise_and_display(sentence, tokenizer):
        """
        Tokenise the sentence using the provided tokenizer and display th
        tokenizer: the tokenizer to use for tokenisation (e.g., GPT-2, BE
        """
        # Tokenise the sentence
        encoding = tokenizer(sentence, return_tensors=None)

        # Generate Tokens, Token IDs and Sequence Length
        token_ids = encoding['input_ids']
        tokens = tokenizer.convert_ids_to_tokens(token_ids)
        seq_length = len(token_ids)

        # Display the results
        print(f"Sentence: {sentence}")
        print(f"Sequence Length: {seq_length}")

        print("-"*40)
        print(f"{'Index':>6}  {'Token':<20}  {'ID':>6}")
        print("-"*40)
        for idx, (token, token_id) in enumerate(zip(tokens, token_ids)):
            print(f"{'idx':>6}  {'token':<20}  {'token_id':>6}")
        print("-"*40)
```

```
In [8]: ## Toeknise and display using DistilGPT2 tokenizer
tokenise_and_display(sentence=sentence, tokenizer=tokenizer1)
```

Sentence: LLMs are powerful tools for natural language understanding.
Sequence Length: 10

Index	Token	ID
0	LL	3069
1	Ms	10128
2	Gare	389
3	Gpowerful	3665
4	Gtools	4899
5	Gfor	329
6	Gnatural	3288
7	Glanguage	3303
8	Gunderstanding	4547
9	.	13

```
In [9]: # Toeknise and display using BERT tokenizer
tokenise_and_display(sentence=sentence, tokenizer=tokenizer2)
```

Sentence: LLMs are powerful tools for natural language understanding.
 Sequence Length: 12

Index	Token	ID
0	[CLS]	101
1	ll	2222
2	##ms	5244
3	are	2024
4	powerful	3928
5	tools	5906
6	for	2005
7	natural	3019
8	language	2653
9	understanding	4824
10	.	1012
11	[SEP]	102

SECTION B

PROMPTS USING DISTILGPT2

PROMPT 1: SUMMARISATION OF PROVIDED TEXT

```
In [10]: paragraph = """True ambition is not merely the pursuit of accolades,
but a methodical dedication to refining one's craft.
Many individuals falter because they equate success with instantaneous re
overlooking the incremental progress necessary for mastery.
A truly ambitious person welcomes failure as a formative experience,
extracting lessons from every setback rather than succumbing to despair.
This tenacious mindset requires a delicate balance of confidence and humi
allowing for continuous adaptation in a rapidly changing environment.
Furthermore, prioritizing long-term goals over immediate gratification di
the visionary from the dreamer.
Ultimately, consistent effort, when paired with strategic thinking, yield
achievement, transforming raw talent into exceptional skill. """
```

```
In [11]: def summarise_with_model(paragraph, model_name):
        """
        Summarise the input paragraph using the specified pre-trained model.
        - model_name: the name of the pre-trained model to use for summarizat
        - the function can be modified to accept additional parameters for ge
        - tokenizer initiated for model_name to allow for flexibility in whic
        """

        ## Load the model and tokenizer, and prepare the input
        tokenizer = AutoTokenizer.from_pretrained(model_name)
        model = AutoModelForSeq2SeqLM.from_pretrained(model_name)

        ## Tokenise the input paragraph
        inputs = tokenizer(paragraph, return_tensors='pt', truncation=True, m

        ## Generate the summary
        summary_ids = model.generate(inputs['input_ids'],
                                     max_length=60,
                                     min_length=20,
```

```

length_penalty=2.0,
num_beams=4,
early_stopping=True
)

## Decode the summary
summary = tokenizer.decode(summary_ids[0], skip_special_tokens=True)

return summary

```

```

In [12]: models = ['sshleifer/distilbart-cnn-12-6', 'facebook/bart-large-cnn']

## loop through models, generate summaries, and store results in a DataFrame
rows = []
for model_name in models:
    summary = summarise_with_model(paragraph, model_name)
    rows.append({
        "model": model_name,
        "summary": summary
    })
df = pd.DataFrame(rows)
display(df)

```

Please make sure the generation config includes `forced\_bos\_token\_id=0`.

Loading weights: 100%|██████████| 358/358 [00:00<00:00, 1268.61it/s, Materializing param=model.shared.weight]

The tied weights mapping and config for this model specifies to tie model.shared.weight to model.decoder.embed_tokens.weight, but both are present in the checkpoints, so we will NOT tie them. You should update the config with `tie\_word\_embeddings=False` to silence this warning

The tied weights mapping and config for this model specifies to tie model.shared.weight to model.encoder.embed_tokens.weight, but both are present in the checkpoints, so we will NOT tie them. You should update the config with `tie\_word\_embeddings=False` to silence this warning

Loading weights: 100%|██████████| 511/511 [00:00<00:00, 1271.20it/s, Materializing param=model.encoder.layers.11.self_attn_layer_norm.weight]

	model	summary
0	sshleifer/distilbart-cnn-12-6	True ambition is not merely the pursuit of ac...
1	facebook/bart-large-cnn	True ambition is not merely the pursuit of acc...

PROMT 2: Q&A FOR FACTUAL QUESTIONS

```

In [13]: qa_roberta = pipeline("question-answering",
                                model="deepset/roberta-base-squad2")

```

Loading weights: 100%|██████████| 199/199 [00:00<00:00, 1740.06it/s, Materializing param=roberta.encoder.layer.11.output.dense.weight]

RobertaForQuestionAnswering LOAD REPORT from: deepset/roberta-base-squad2

Key	Status	
roberta.embeddings.position_ids	UNEXPECTED	

Notes:

- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.

```
In [14]: qa_distilbert = pipeline("question-answering",
                                model="distilbert-base-uncased-distilled-squad")
```

```
Loading weights: 100%|██████████| 102/102 [00:00<00:00, 886.90it/s, Materializing param=qa_outputs.weight]
```

```
In [15]: ## The context for question answering – can be modified to use a different
context = """
Large Language Models (LLMs) are advanced artificial intelligence systems
on vast amounts of text data. They are designed to understand and generate
human-like text across a wide range of tasks, including translation, summarization,
question answering, and creative writing.

LLMs such as GPT, BERT, and LLaMA are built using transformer architectures
which allow them to capture long-range dependencies in language. They are
trained on large corpora and can be fine-tuned for specific downstream tasks.

Despite their capabilities, LLMs have notable limitations. They can produce
factually incorrect outputs, known as hallucinations, and may reflect biases
present in their training data. They also require significant computational
resources to train and deploy.

Researchers and engineers continue to improve LLMs by developing more efficient
training techniques, better alignment methods, and robust evaluation benchmarks
to ensure safe and reliable use in real-world applications.
"""
```

```
In [16]: def generate_answers(question, context, qa):
        """
        Generate an answer to the question based on the provided context using
        the question-answering pipeline.
        - question: the question to be answered
        - context: the context from which the answer should be derived
        - qa: the question-answering pipeline to use for generating the answer
        """
        result = qa(question=question, context=context)
        return result['answer']
```

```
In [17]: questions = [
        "What are LLMs trained on?",
        "What architecture do LLMs use?",
        "What are the limitations of LLMs?",
        "What is a hallucination in the context of LLMs?",
        "How can LLMs be improved?",
    ]

## Loop through questions, generate answers using both models, and display them
for question in questions:
    answer_roberta = generate_answers(question, context, qa_roberta)
    answer_distilbert = generate_answers(question, context, qa_distilbert)

    print(f"Question: {question}")
    print(f"Answer (RoBERTa): {answer_roberta}")
    print(f"Answer (DistilBERT): {answer_distilbert}")
    print("-"*50)
```


Question: What are LLMs trained on?

Answer (RoBERTa): vast amounts of text data

Answer (DistilBERT): large corpora

Question: What architecture do LLMs use?

Answer (RoBERTa): transformer architectures

Answer (DistilBERT): transformer architectures

Question: What are the limitations of LLMs?

Answer (RoBERTa): They can produce factually incorrect outputs

Answer (DistilBERT): They can produce factually incorrect outputs

Question: What is a hallucination in the context of LLMs?

Answer (RoBERTa): incorrect outputs

Answer (DistilBERT): factually incorrect outputs

Question: How can LLMs be improved?

Answer (RoBERTa): more efficient training techniques, better alignment methods, and robust evaluation benchmarks

Answer (DistilBERT): more efficient training techniques, better alignment methods, and robust evaluation benchmarks

PROMPT 3: GENERATE CREATIVE TEXT

```
In [18]: # generator already initialised above
def generate_text(generator, prompt):
    """
    Generate text using the specified generator and prompt, with parameters:
    - generator: the text generation pipeline to use (e.g., distilgpt2)
    - temperature: controls the randomness of the generated text (higher is more random)
    - top_p: controls the diversity of the generated text by limiting the cumulative probability
    - do_sample: whether to use sampling for text generation (if False, uses greedy decoding)
    - repetition_penalty: penalizes repeated tokens to encourage more diverse outputs
    - num_return_sequences: the number of different sequences to generate
    """
    generated = generator(prompt,
                          max_new_tokens=100,
                          max_length=50,
                          num_return_sequences=3,
                          temperature=1.0, # should generate same output
                          top_p=0.9,
                          do_sample=True,
                          repetition_penalty=1.2)
    return [r['generated_text'] for r in generated]
```

```
In [19]: prompts = [
    "The future of AI is ",
    "Language flows like water through the veins of human communication, ",
    "The year is 2045. Machines have learned to speak, but humans have fo
]

## Loop through prompts, generate text, and store results in a DataFrame
## Uses the generator for distilgpt2, but can be modified to use a differ
```

```

rows = []
for prompt in prompts:
    text = generate_text(generator, prompt)
    for idx, gen in enumerate(text,1):
        rows.append({
            "prompt": prompt,
            "model": "distilgpt2",
            "Generated Response #": idx,
            "generated_text": gen
        })

import pandas as pd
df = pd.DataFrame(rows)
pd.set_option("display.max_colwidth", None)
display(df)

```

Passing `generation\_config` together with generation-related arguments= ({'num_return_sequences', 'max_length', 'temperature', 'max_new_tokens', 'top_p', 'do_sample', 'repetition_penalty'}) is deprecated and will be removed in future versions. Please pass either a `generation\_config` object OR all generation parameters explicitly, but not both.

Setting `pad\_token\_id` to `eos\_token\_id`:50256 for open-end generation. Both `max\_new\_tokens` (=100) and `max\_length` (=50) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Setting `pad\_token\_id` to `eos\_token\_id`:50256 for open-end generation. Both `max\_new\_tokens` (=100) and `max\_length` (=50) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Setting `pad\_token\_id` to `eos\_token\_id`:50256 for open-end generation. Both `max\_new\_tokens` (=100) and `max\_length` (=50) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

	prompt	model	Generated Response #	generated_
0	The future of AI is	distilgpt2	1	The future of AI is ills and the most beautiful thing that can happen. The good news here at Machine Intelligence is that it could be improved, but we don't have the time right now for a project like this with this technology in our ha
1	The future of AI is	distilgpt2	2	The future of AI is already a serious debate and it has the potential to change very rapidly. At first glance there are some ideas out in this forum but for too much discussion we shall keep exploring them as they come up again so you can see where those theories take us next... In all cases however my advice would be to think about what I am saying here goes into depth - because it is our most fundamental concepts have been tested and thoroughly now then why didn't other people realize this yet? That at least give
2	The future of AI is	distilgpt2	3	The future of AI is _____
3	Language flows like water through the veins of human communication, shaping our thoughts and connecting us across time and space.	distilgpt2	1	Language flows like water through the veins of human communication, shaping our thoughts and connecting us across time and space. This is all about a person's awareness that we exist in real life rather than merely as being connected to it - by giving us information over more direct transmission via internet or digital channels within one context so they can be accessed without having ever looked at any particular subject for such access (e-mail). And this means that an individual will feel when he hears his own words. The meaning of my word(s). There are no other ways your brain translates into people living there
4	Language flows like water through the veins of human communication, shaping our thoughts and connecting us across time and space.	distilgpt2	2	Language flows like water through the veins of human communication, shaping our thoughts and connecting us across time and space. The notion that we can learn from this in order to move past obstacles doesn't always fit with my worldview or experience all as much. I don't feel compelled by a lack of evidence on social media when I am forced into thinking about how best for me to succeed (or just not) is learning a little it often gets lost between myself and others which makes some people hesitate over what to have learned along those lines: "How do you know your job has
5	Language flows like water through the veins of human communication, shaping our thoughts and connecting us across time and space.	distilgpt2	3	Language flows like water through the veins of human communication, shaping our thoughts and connecting us across time and space. A group in Singapore celebrates their achievements by giving a shout out to each other with an emoji during which they express hope for harmony within society — despite political divisions over this issue being raised at different levels or discussed often online on social media (Facebook). While there may not seem much that could be done about it now though, some are calling for attention both as well but especially since so many people also make fun around such issues related specifically to immigration policy: -I have seen

	prompt	model	Generated Response #	generated_
6	The year is 2045. Machines have learned to speak, but humans have forgotten to listen.	distilgpt2	1	The year is 2045. Machines have learned to speak but humans have forgotten to listen. \nI am an engineer working on the project for a local company called Avant-Systems and my work will not just be open until it is completed next week! The world won't know more about this stuff then I would like to share
7	The year is 2045. Machines have learned to speak, but humans have forgotten to listen.	distilgpt2	2	The year is 2045. Machines have learned to speak but humans have forgotten to listen. ***** August of that same month, Google announced the creation of an "intelligence company," founded by John O'Brien's cofounder and chief technology officer, Steve Ballmer. The firm launched its first services platform in late April on behalf as well as Apple Maps for iOS — it has not been mentioned until now so far; we are waiting eagerly for details about what such a future might look like (hopefully). And then here comes Facebook CEO Mark Zuckerberg who reportedly p
8	The year is 2045. Machines have learned to speak, but humans have forgotten to listen.	distilgpt2	3	The year is 2045. Machines have learned to speak but humans have forgotten to listen.

SECTION C: EMBEDDINGS WITH GENSIM

In [20]: `pip install gensim`

```
Requirement already satisfied: gensim in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (4.4.0)
Requirement already satisfied: numpy>=1.18.5 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from gensim) (2.4.2)
Requirement already satisfied: scipy>=1.7.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from gensim) (1.17.0)
Requirement already satisfied: smart_open>=1.8.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from gensim) (7.5.1)
Requirement already satisfied: wrapt in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (from smart_open>=1.8.1->gensim) (2.1.1)
```

[notice] A new release of pip is available: 26.0 -> 26.0.1

[notice] To update, run: `pip install --upgrade pip`

Note: you may need to restart the kernel to use updated packages.

DOWNLOAD GLOVE

- ONE TIME ONLY...
- SECOND ATTEMPT WILL PRINT FILE ALREADY EXISTS! SKIPPING DOWNLOAD

```
In [21]: glove_txt = "glove.6B.50d.txt"
glove_zip = "glove.6B.zip"
glove_url = "https://nlp.stanford.edu/data/glove.6B.zip"

if not os.path.exists(glove_txt):
    print("Downloading GloVe (≈ 862 MB, one-time)...")

    # Bypass SSL certificate verification (safe for known URLs on trusted
    ssl_ctx = ssl.create_default_context()
    ssl_ctx.check_hostname = False
    ssl_ctx.verify_mode = ssl.CERT_NONE

    opener = urllib.request.build_opener(
        urllib.request.HTTPSHandler(context=ssl_ctx)
    )
    urllib.request.install_opener(opener)

    urllib.request.urlretrieve(glove_url, glove_zip)

    with zipfile.ZipFile(glove_zip, "r") as z:
        z.extract(glove_txt)
    os.remove(glove_zip)
    print("Done.")
else:
    print("GloVe file already exists. Skipping download.")
```

GloVe file already exists. Skipping download.

LOAD GLOVE IN MEMORY

```
In [22]: # Load GloVe embeddings (50-dimensional, trained on Wikipedia + Gigaword)
glove = api.load("glove-wiki-gigaword-50")
```

EXTRACT TOP 10 VECTORS FOR SELECTED WORDS AND SIMILAR WORDS TO SELECTED WORDS

```
In [23]: words = ["king", "queen", "diamond"]

for word in words:
    vector = glove[word]
    similar = glove.most_similar(word, topn=5)

    display(Markdown(f"### `{word}`"))

    # First 10 vector values
    print("First 10 vector values:")
    vec_df = pd.DataFrame([vector[:10]], columns=[f"d{i}" for i in range(
    display(vec_df.round(4))

    # Top 5 similar words
    print("Top 5 most similar words:")
    sim_df = pd.DataFrame(similar, columns=["Word", "Similarity Score"])
    sim_df["Similarity Score"] = sim_df["Similarity Score"].round(4)
```

```
display(sim_df)
print()
```

king

First 10 vector values:

	d0	d1	d2	d3	d4	d5	d6	d7	d8	d9
0	0.5045	0.6861	-0.5952	-0.0228	0.6005	-0.135	-0.0881	0.4738	-0.618	-0.3101

Top 5 most similar words:

	Word	Similarity Score
0	prince	0.8236
1	queen	0.7839
2	ii	0.7746
3	emperor	0.7736
4	son	0.7667

queen

First 10 vector values:

	d0	d1	d2	d3	d4	d5	d6	d7	d8	d9
0	0.3785	1.8233	-1.2648	-0.1043	0.3583	0.6003	-0.1754	0.8377	-0.0568	-0.758

Top 5 most similar words:

	Word	Similarity Score
0	princess	0.8515
1	lady	0.8051
2	elizabeth	0.7873
3	king	0.7839
4	prince	0.7822

diamond

First 10 vector values:

	d0	d1	d2	d3	d4	d5	d6	d7	d8	d9
0	-0.4958	0.7842	-0.606	1.3967	0.2889	-0.2058	-0.1074	-0.3325	1.3608	0.1509

Top 5 most similar words:

	Word	Similarity Score
0	gold	0.7715
1	diamonds	0.7663
2	gem	0.7375
3	silver	0.7210
4	jewel	0.7102

SENTENCE LEVEL REPRESENTATION AND SIMILARITY - CONTEXT

- DETERMINE AVERAGE WORD VECTOR REPRESENTING A SENTENCE FROM A GROUP OF SENTENCES
- DETERMINE SIMILARITY BETWEEN SENTENCES (COSINE)

```
In [24]: sentences = [
    "What are LLMs trained on?",
    "What architecture do LLMs use?",
    "What are the limitations of LLMs?",
    "What is a hallucination in the context of LLMs?",
    "How can LLMs be improved?",
]
```

```
In [25]: ## Compute average word vector per sentence
def sentence_vector(sentence, model):
    """
    Average the GloVe vectors of all known words in the sentence.
    """
    tokens = sentence.lower().replace("?", "").split() ## since all sente
    vectors = [model[w] for w in tokens if w in model]
    return np.mean(vectors, axis=0) if vectors else np.zeros(model.vector

## Each sentence is represented as a 50-dimensional vector (average of wo
sent_vectors = np.array([sentence_vector(s, glove) for s in sentences])
```


```
In [26]: ## Cosine similarity matrix between sentence vectors
sim_matrix = cosine_similarity(sent_vectors)

## Short labels for readability
labels = [f"S{i+1}" for i in range(len(sentences))]

sim_df = pd.DataFrame(sim_matrix, index=labels, columns=labels).round(3)

## Heatmap
fig, ax = plt.subplots(figsize=(7, 5))
sns.heatmap(
    sim_df, annot=True, fmt=".2f", cmap="YlOrRd",
    vmin=0, vmax=1, linewidths=0.5, ax=ax
)
ax.set_title("Sentence Similarity Matrix (GloVe avg vectors)", pad=12)
plt.tight_layout()
plt.show()
```

```
## Legend – which label maps to which sentence
print("\nSentence Legend:")
for label, sentence in zip(labels, sentences):
    print(f" {label}: {sentence}")
```



Sentence Legend:

S1: What are LLMs trained on?
 S2: What architecture do LLMs use?
 S3: What are the limitations of LLMs?
 S4: What is a hallucination in the context of LLMs?
 S5: How can LLMs be improved?

SECTION D: SENTIMENT ANALYSIS

```
In [27]: ## LOAD MODEL
# BERT fine-tuned on SST-2 (Stanford Sentiment Treebank – movie reviews)
# Binary output: POSITIVE / NEGATIVE + confidence score
classifier = pipeline("sentiment-analysis",
                      model="textattack/bert-base-uncased-SST-2")
```

Loading weights: 100%|██████████| 201/201 [00:00<00:00, 240.78it/s, Materializing param=classifier.weight]

```
In [28]: ## SAMPLE DATASET
## Includes movie reviews (original training data – like samples)
## Also includes restaurant and product reviews
reviews = [
    ## Film reviews
    "An absolute masterpiece. The performances were outstanding and the s
    "A decent watch. Nothing groundbreaking but enjoyable enough for a Su
    "Oh brilliant – another two hours of my life I'll never get back. Tru
    "Dull, predictable, and far too long. A complete waste of time.", # S
```



```

    "The film runs 142 minutes and is based on the 2003 novel of the same
    "The visuals were stunning and the score was beautiful, but the plot
    "A slow, confusing first hour – but the finale was so powerful it red

    # Domain shift: product review
    "Battery life is excellent and the screen is gorgeous. Best phone I've
    # Domain shift: restaurant review
    "Overpriced for what it is. Tiny portions and the service was dismiss
    ## Could be movie, restaurant, or a consumable product
    "It had potential but never really delivered. Forgettable at best.",
]

```

```

In [29]: ## Run pre-built classifier on a new set of reviews
results = classifier(reviews)

label_map = {"LABEL_0": "NEGATIVE", "LABEL_1": "POSITIVE"}
rows = []
idx = 0
for review, result in zip(reviews, results):
    rows.append({
        "Index"      : idx,
        "Review"      : review,
        "Sentiment"   : label_map[result["label"]],
        "Confidence"  : round(result["score"], 4)
    })
    idx += 1

df = pd.DataFrame(rows)
pd.set_option("display.max_colwidth", None)

```

```

In [30]: def color_row(row):
    color = "#d4edda" if row["Sentiment"] == "POSITIVE" else "#f8d7da"
    return [f"background-color: {color}"] * len(row)

styled_df = df[["Index", "Review", "Sentiment", "Confidence"]].style.apply
display(styled_df)

```

	Index	Review	Sentiment	Confidence
0	0	An absolute masterpiece. The performances were outstanding and the story kept me hooked.	POSITIVE	0.999700
1	1	A decent watch. Nothing groundbreaking but enjoyable enough for a Sunday afternoon.	POSITIVE	0.999200
2	2	Oh brilliant — another two hours of my life I'll never get back. Truly unforgettable.	POSITIVE	0.999600
3	3	Dull, predictable, and far too long. A complete waste of time.	NEGATIVE	0.999200
4	4	The film runs 142 minutes and is based on the 2003 novel of the same name.	POSITIVE	0.852000
5	5	The visuals were stunning and the score was beautiful, but the plot was an incoherent mess.	NEGATIVE	0.989400
6	6	A slow, confusing first hour — but the finale was so powerful it redeemed everything.	POSITIVE	0.999300
7	7	Battery life is excellent and the screen is gorgeous. Best phone I've owned.	POSITIVE	0.999700
8	8	Overpriced for what it is. Tiny portions and the service was dismissive.	NEGATIVE	0.998500
9	9	It had potential but never really delivered. Forgettable at best.	NEGATIVE	0.997000

HOW CAN THIS BE USED IN BUSINESS WORKFLOW, WHAT COULD BE THE LIMITATIONS?

```
In [32]: display(Markdown("""
### Place in Business Workflow Reflection

BERT based sentiment classifier demonstrates strong performance on clear-

At least three important limitations are observed:
- sarcasm in Review 2 was misclassified based on positive words
- Model struggles at evaluating factual statement
  - Neutral review in Review 4 (length of film) was construed as positive
- Product reviews (Review 7 & 8) while correctly classified, are based on

For professional deployment, these findings suggest that BERT is a strong
But for high-stakes decisions — such as escalating a customer complaint o
"""))
```

Place in Business Workflow Reflection

BERT based sentiment classifier demonstrates strong performance on clear-cut reviews, but struggles with more nuanced (mixed-tone) or switch in application / domain (trained on movie reviews applied on product and restaurant reviews).

At least three important limitations are observed:

- sarcasm in Review 2 was misclassified based on positive words
- Model struggles at evaluating factual statement
 - Neutral review in Review 4 (length of film) was construed as positive *albeit* with low confidence
- Product reviews (Review 7 & 8) while correctly classified, are based on keywords rather than true context

For professional deployment, these findings suggest that BERT is a strong **first-pass** filter for high-volume feedback. But for high-stakes decisions — such as escalating a customer complaint or flagging brand risk — require more mature and context-aware approaches like RAG pipelines, domain-specific fine-tuning, and human-in-the-loop review for edge cases.

In []: